

## PART I – Background and Properties

### Four Parts of Cassandra Tutorial

- History and Background
- Installation
- Speed Contest Vs MongoDB
- Generate A Speed Layer

### History and Background

- Apache Cassandra - An open source, distributed /decentralized database, provides highly available service with no single point of failure.
- Notable Features:
  - **Elastic scalability** - Cassandra is highly scalable; it allows to add more hardware to accommodate more customers and more data as per requirement.
  - **Always on architecture** - Cassandra has no single point of failure and it is continuously available for business-critical applications that cannot afford a failure.
  - **Fast linear-scale performance** - Cassandra is linearly scalable. Therefore, it maintains a quick response time.
  - **Flexible data storage** - Dynamically accommodate structured, semi-structured, and unstructured according to need.
  - **Easy data distribution** - Cassandra provides the flexibility to distribute data where you need by replicating data across multiple data centers.
  - **Transaction support** - Cassandra supports properties like Atomicity, Consistency, Isolation, and Durability (ACID).
  - **Fast writes** - run on cheap commodity hardware. It performs blazingly fast writes and can store hundreds of terabytes of data, without sacrificing the read efficiency.

### Pros:

- It is scalable, fault-tolerant, and consistent.
- It is a column-oriented database.
- Its distribution design is based on Amazon's Dynamo and its data model on Google's Bigtable.
- Created at Facebook, it differs sharply from relational database management systems.
- Cassandra implements a Dynamo-style replication model with no single point of failure, but adds a more powerful "column family" data model.
- Cassandra is being used by some of the biggest companies such as Facebook, Twitter, Cisco, Rackspace, ebay, Twitter, Netflix, and more.

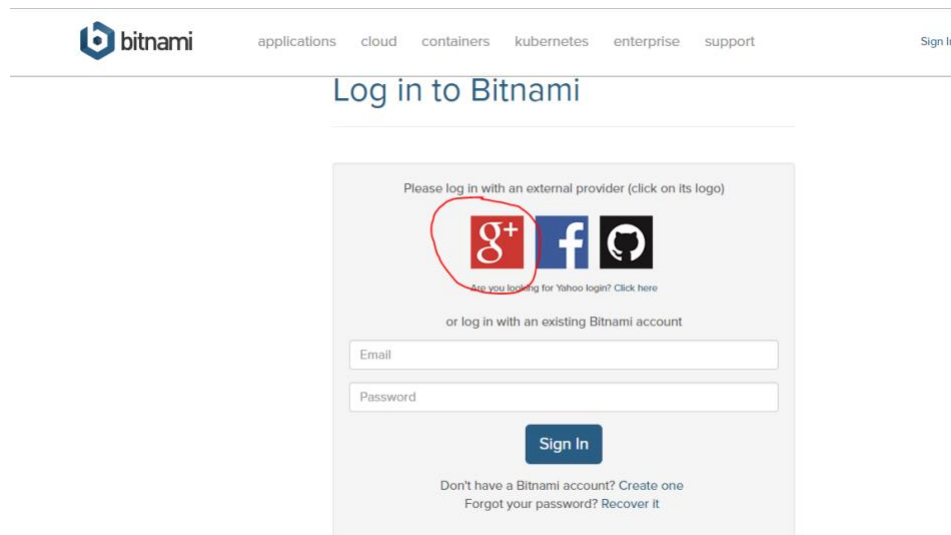
## PART II – Installation

### Cassandra Cloud Set Up Steps

- Create google account

- Open <https://cloud.google.com>
- Sign up for free trial
- Create a project on the google cloud console
- Open “<https://google.bitnami.com>” and sign in with Google
- Create a virtual machine

1. Go to the URL [https://bitnami.com/sign\\_in](https://bitnami.com/sign_in) and login with your own google account:



2.
  - a). Type in the name you want to name your Cassandra vm (Step 1)
  - b). Select Cassandra from the list of virtual machine instances (Step 2)
  - c). Select the google cloud project (project we just created) (Step 3)
  - d) Select create to create the project (Step 4)

Cloud Platform

1

NAME

my-cassandra-server

2

IMAGE

Cassandra v3.11.1-4 (Debian 8)

3

CLOUD ACCOUNT

Cassandra speed layer (cassandra-speed-layer)

NETWORK

default

DISK TYPE

☐ Solid State
☒ Magnetic Disk

DISK SIZE

10 GB

\$0.40 /mo

SERVER SIZE

☐ n1-highcpu-2  
(\$36.23 /mo) \$0.071 /hr

☒ n1-standard-2  
(\$48.55 /mo) \$0.095 /hr

☐ n1-highmem-2  
(\$60.50 /mo) \$0.118 /hr

REGION

us-central1-f


Estimated Monthly cost: \$48.95


CANCEL

4 CREATE


3. The Cloud will give a default test cluster, you can also create your own clusters and name it as your personal need:

## New Project

 You have 9 projects remaining in your quota. [Learn more.](#)

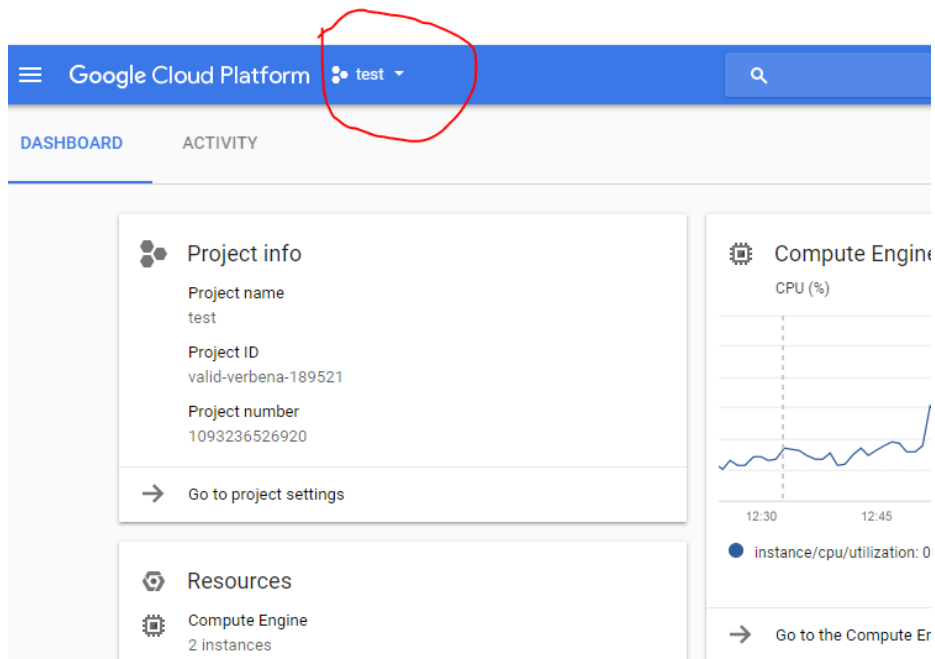
Project name 

My Project 82071


Your project ID will be ~~sunlit-pipe-189618~~  [Edit](#)

Create


Cancel



The screenshot shows the Google Cloud Platform dashboard for a project named 'test'. The top navigation bar is blue with the Google Cloud logo, the project name 'test', and a search icon. Below the navigation bar, there are two tabs: 'DASHBOARD' and 'ACTIVITY'. The main content area is divided into two columns. The left column contains two sections: 'Project info' and 'Resources'. The 'Project info' section displays the project name 'test', project ID 'valid-verbena-189521', and project number '1093236526920', with a link to 'Go to project settings'. The 'Resources' section shows 'Compute Engine' with '2 instances'. The right column features a 'Compute Engine' section with a line graph showing CPU usage over time, with a legend indicating 'instance/cpu/utilization: 0' and a link to 'Go to the Compute Engine page'.

Google Cloud Platform  test

DASHBOARD ACTIVITY


 **Project info**


Project name  
test


Project ID  
valid-verbena-189521

Project number  
1093236526920

→ Go to project settings

 **Resources**

 Compute Engine  
2 instances

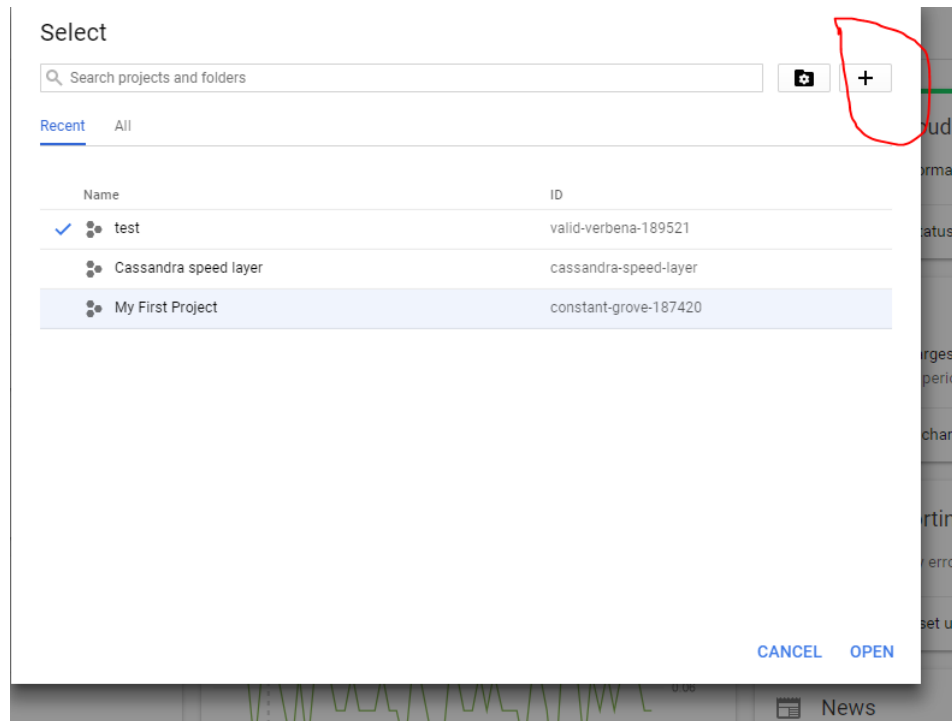
 **Compute Engine**

CPU (%)

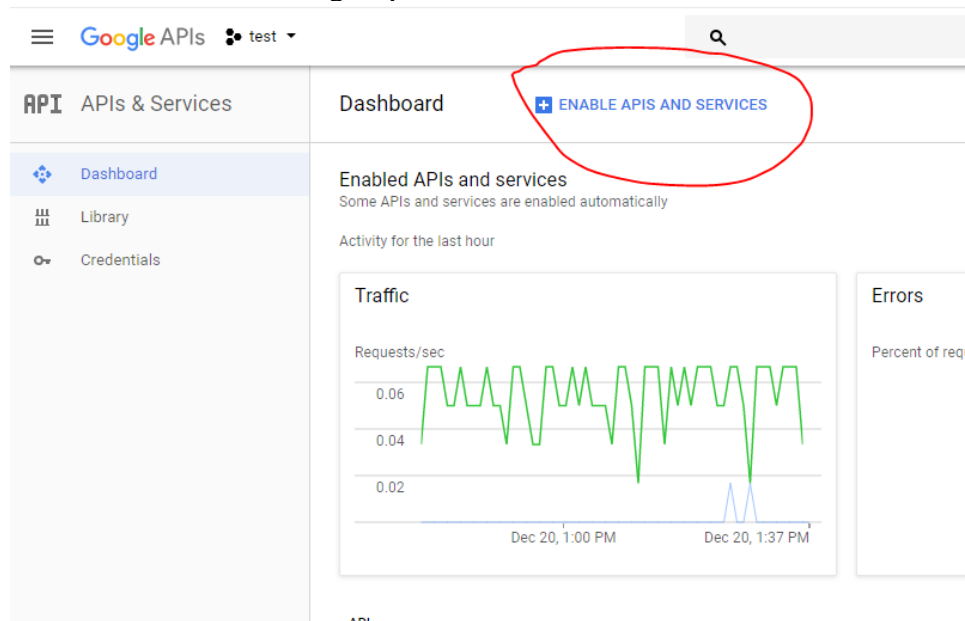
12:30 12:45

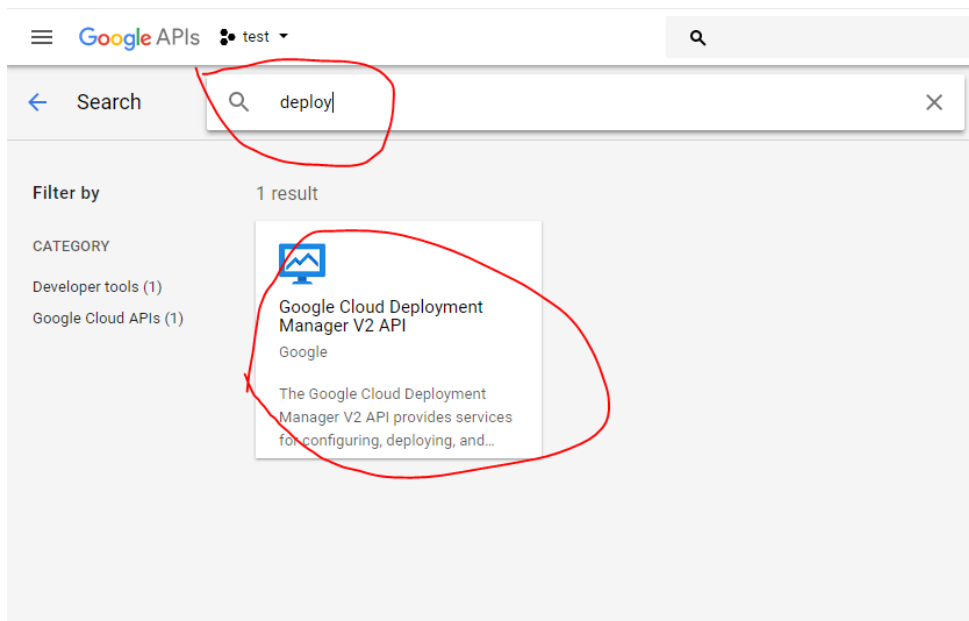
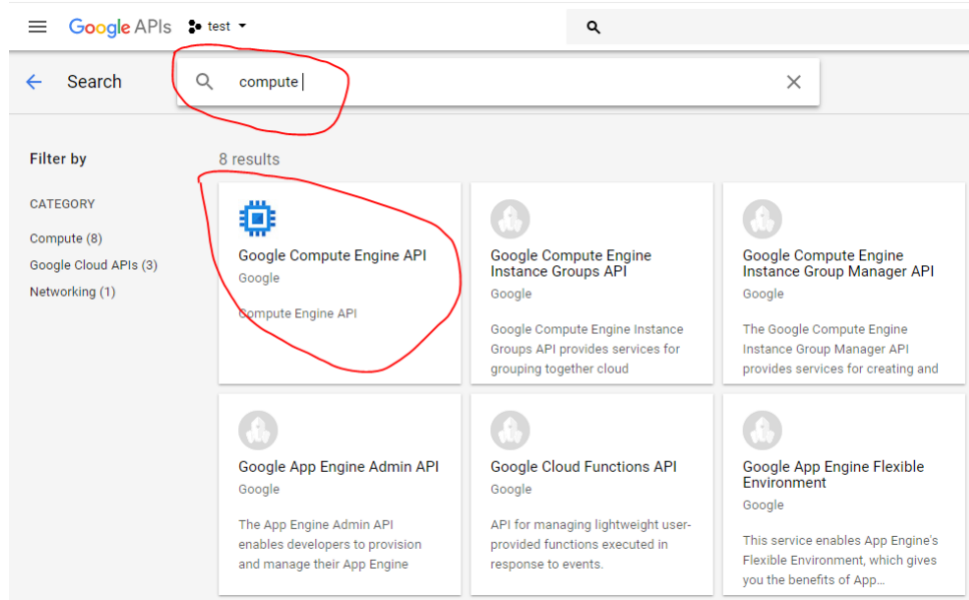
● instance/cpu/utilization: 0

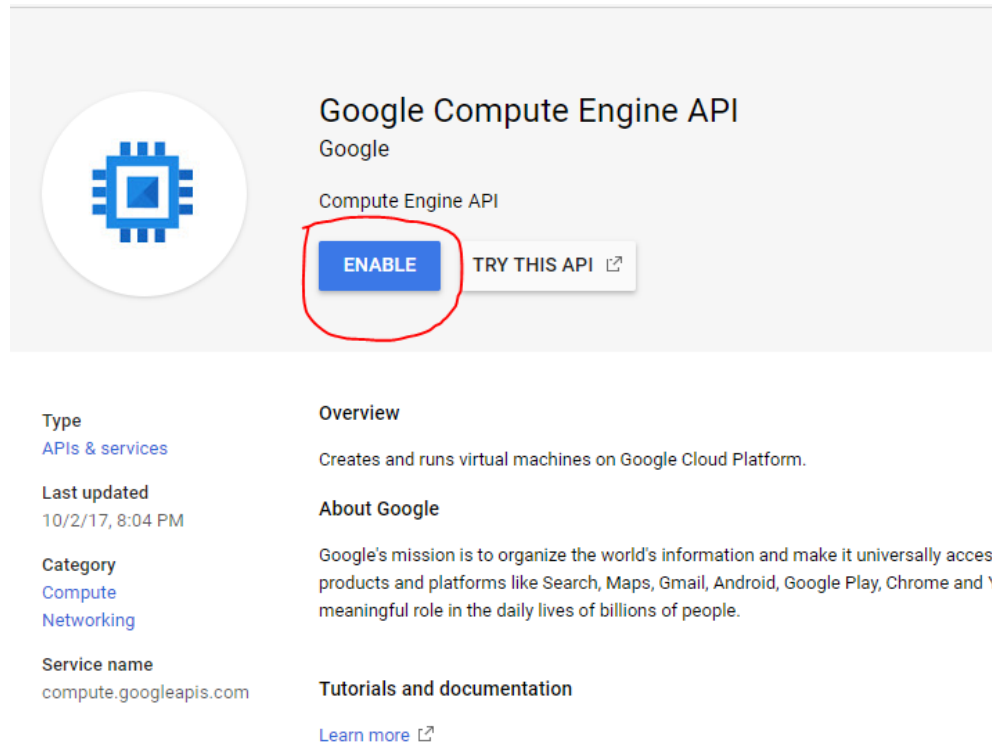
→ Go to the Compute Engine page



4. Enable APIs and Services as following steps:







**Google Compute Engine API**  
Google

Compute Engine API

**ENABLE** TRY THIS API ↗

**Type**  
[APIs & services](#)

**Last updated**  
10/2/17, 8:04 PM

**Category**  
[Compute](#)  
[Networking](#)

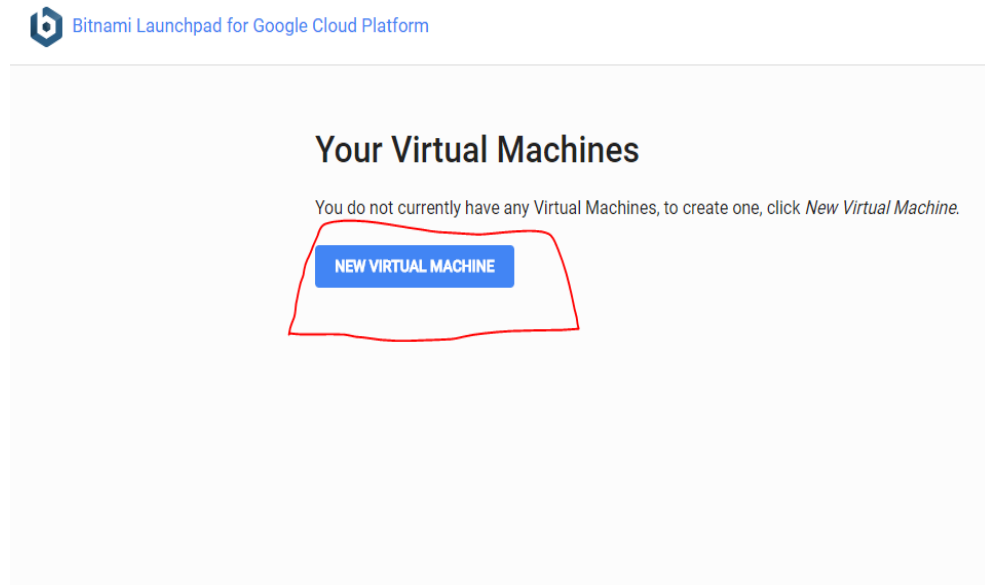
**Service name**  
compute.googleapis.com

**Overview**  
Creates and runs virtual machines on Google Cloud Platform.

**About Google**  
Google's mission is to organize the world's information and make it universally access products and platforms like Search, Maps, Gmail, Android, Google Play, Chrome and Y meaningful role in the daily lives of billions of people.

**Tutorials and documentation**  
[Learn more](#) ↗

5. Once logged in accept the license agreement. Follow the image step below:



**Bitnami Launchpad for Google Cloud Platform**

## Your Virtual Machines

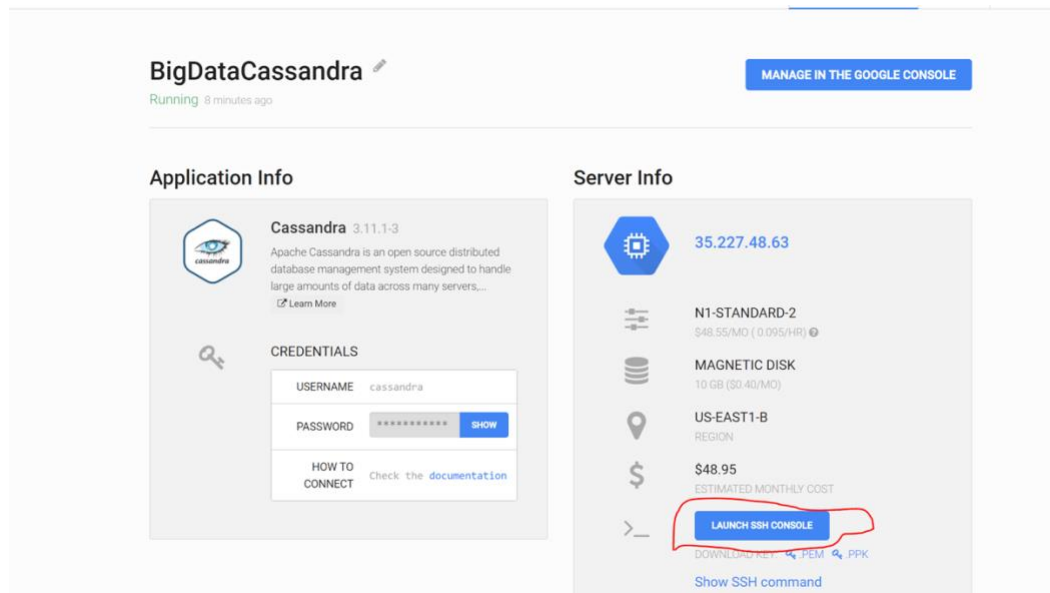
You do not currently have any Virtual Machines, to create one, click *New Virtual Machine*.

**NEW VIRTUAL MACHINE**

6. Multiple node cluster:

1). Go to the bitnami console using "<https://google.bitnami.com>"

- 2). Sign in and select virtual machines from the top right of the screen
- 3). Select the cassandra virtual machine you have setup and the screen below will show up\



6. In the online console, follow these steps:

- 1). Download the personal key from Bitnami.
- 2). Open local terminal, change permission to public by type in:  
`chmod 600 bitnami-google-cassandra-speed-layer.pem.`
- 3). Run: `ssh -i bitnami-google-cassandra-speed-layer.pem`
- 4). Connect to cluster:  
`cqlsh -u cassandra 23.251.158.183 -p N4pKBV8AjKiU`
- 5). You can change password to yours:

`cqlsh> ALTER USER cassandra with PASSWORD 'NEWPASSWORD';`

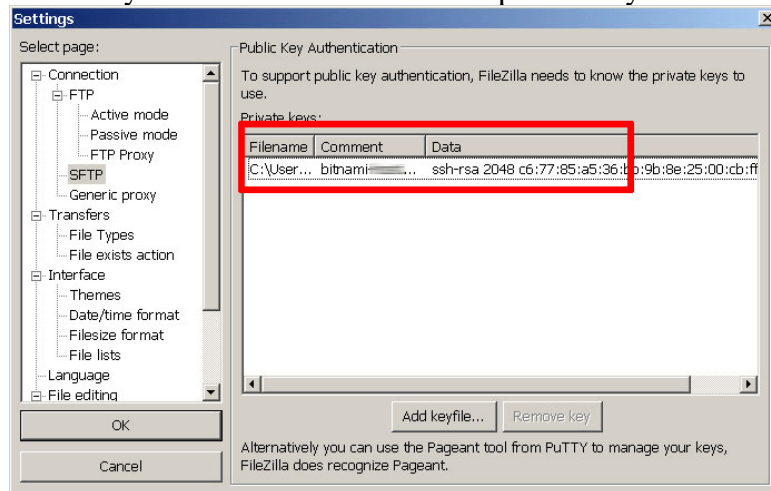
6). Restart, all set!

**Get connected to the FileZilla** (To upload local files to the server with SFTP.)

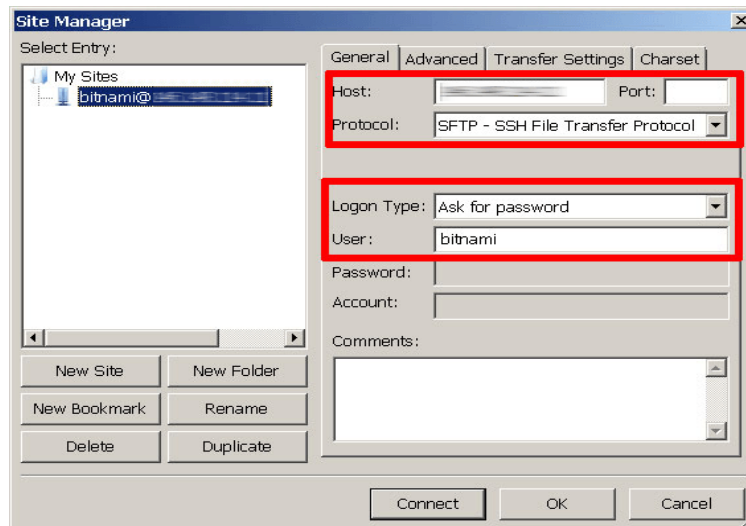
- Download and install FileZilla, link: <https://filezilla-project.org/>
- Launch FileZilla and use the "Edit -> Settings" command to bring up FileZilla's configuration settings.



- Use the "Add keyfile" command to select the private key file for the server:



- Use the "File -> Site Manager -> New Site" command to bring up the FileZilla Site Manager.
- Enter your server host name and specify *bitnami* as the user name.
- Select "SFTP" as the protocol and "Ask for password" as the logon type.
- Use the "Connect" button to connect to the server and begin an SFTP session. You might need to accept the server key, by clicking "Yes" or "OK" to proceed.



## Generate Speed Views


### Nodes and Clusters

- Every node is a peer
- Ring structure (not master/slave)
- Tokens and hashing

- Elastic scaling (remove or add node from (to) running cluster, upgrade when it's running, ...)

### Writes (what happens within each node)

- Data is first written to a commit log for durability. Your data is safe in Cassandra
- Then written to a memtable in memory
- Once the memtable becomes full, it is flushed to an SSTable (sorted strings table)
- Writes are atomic at the row level; all columns are written or updated, or none are. RDBMS-styled transactions are not supported



Cassandra is known for being the fastest database in the industry where write operations are concerned.

### Consistency

- Every node acts as a coordinator
- Tunable consistency levels
- Strong consistency ( $R + W > N$ )
- Eventually consistency ( $R + W < N$ )

### Cassandra Data Model

- Query language (CQL) looks like SQL
- Data model is tabular (like relational database)
  - Tall, but narrow (6-10 columns)
  - Scalable → can keep millions or billions of rows
- Differences:
  - Cassandra do not do JOIN tables
  - Denormalize a lot (at scale we do not JOIN, it slows the read down)
  - Keyspace is namespace (container) for tables

### Cassandra Keys

- Record ID:
  - Primary key:
  - Uuid:
    - surrogate primary key
    - why not just integer?
- Keys:

- Primary keys uniquely identify rows (like RB)
- Each key has 2 parts:
  - Partition key: group of rows in the table is guaranteed is located in the same node
  - Clustering key: provide ordering to the rows in the table
- You need to put a column in a key so that it can participate in the query

## Connected to Clusters:

```
auth_provider = PlainTextAuthProvider(username='cassandra',
                                       password='R7vCBaqL84yB')

contact_points = ['35.227.48.63', '35.227.96.7', '35.196.169.200']

cluster = Cluster(
    contact_points=contact_points, auth_provider=auth_provider,
    load_balancing_policy= TokenAwarePolicy(DCAwareRoundRobinPolicy(local_dc='us-east1')),
    default_retry_policy = RetryPolicy()
)

# connect to the cluster
session = cluster.connect()
cassandraprojectbigdata@bitnami-cassandra-dm-6397:~$ /opt/bitnami/cassandra/bin/nodetool status
Datacenter: us-east1
=====
Status=Up/Down
-- State=Normal/Leaving/Joining/Moving
-- Address      Load       Tokens     Owns    Host ID                               Rack
UN 35.227.96.7    335.85 KiB 256        ?       68b11949-7590-42f9-ab71-3cc831c25ce0  b
UN 35.227.48.63  337.57 KiB 256        ?       d72d7f76-6707-4598-b917-47af85156910  b
UN 35.196.169.200 351.14 KiB 256        ?       47b9fd66-98c5-45e9-9aac-b389a5f4accf  b
Note: Non-system keyspaces don't have the same replication settings, effective ownership information is meaningless
```

## To Create Tables:

```
KEYSPACE = 'demo'
session.execute('DROP KEYSPACE IF EXISTS %s' % KEYSPACE)
session.execute("""
    CREATE KEYSPACE IF NOT EXISTS %s
    WITH replication = { 'class': 'SimpleStrategy', 'replication_factor': '2' }
    AND durable_writes = true;
    """ % KEYSPACE)

class movies_by_actor(Model):
    actor = c.Text(primary_key=True)
    release_year = c.Integer(primary_key=True, clustering_order="DESC")
    movie_id = c.UUID(primary_key=True, clustering_order="ASC")
    genres = c.Set(c.Text())
    rating = c.Float()
    title = c.Text()
```

```
cassandra@cqlsh> select * from demo.movies_by_actor;
```

actor	release_year	movie_id	genres	rating	title
Tom Hanks	2016	d69fba92-3546-40cb-b7b2-8493017053ca	['biography', 'drama']	7.5	Sully
Tom Hanks	2015	09e25d48-6a87-4df8-979f-9f5dda016d45	['drama', 'thriller']	7.6	Bridge of Spies
Tom Hanks	1994	f235b199-0880-448f-9015-5a25ea7671f7	['comedy', 'drama']	8.8	Forrest Gump
Emma Stone	2016	7e12c85b-e57d-4f77-96fc-9fb1e0277aba	['drama', 'romance']	8.1	La La Land

## Speed Test: Cassandra vs MongoDB

- According to the Cassandra website (Apache, 2016): “Cassandra consistently outperforms popular NoSQL alternatives in benchmarks and real applications”
- Datastax (Datastax, 2017): “For mixed operational and analytic workloads typical to modern Web, Mobile and IOT applications, Cassandra performed six times faster than HBase and 195 times faster than MongoDB.”

## Test: Cassandra vs MongoDB – Procedure

- Cassandra and MongoDB node (one each) set up on Google Cloud Platform
- Three trials: In each, create 1000 pretend movies (fictitious films that have the same fields as that from the movie database) and insert them into the respective database.
- Trials differ in the amount of times we repeat the process (1, 100, 1000).
- Statistics collected displaying the time (in seconds)
  - Total time for # of inserts
  - Average time for # of inserts

## Speed Test: Cassandra vs MongoDB – Procedure

- Cassandra and MongoDB node (one each) set up on Google Cloud Platform
- Two trials: Create and insert 100 or 1000 movies.
- Statistics collected displaying the time (in seconds)
- Each also had some specific tests (1 movie insert for MongoDB, adding movies by actor for Cassandra; see GitHub for full results)

## References:

1. Apache, 2016. <http://cassandra.apache.org/>
2. Datastax, 2017. <https://www.datastax.com/apache-cassandra-leads-nosql-benchmark>
3. Arunkumar U, June 2017. [https://medium.com/@arun\\_74827/what-is-apache-cassandra-what-are-the-features-of-it-a4b26b860d07](https://medium.com/@arun_74827/what-is-apache-cassandra-what-are-the-features-of-it-a4b26b860d07)
4. Bitnami Docs, <https://docs.bitnami.com/google/infrastructure/cassandra/#description>